

Small measured benefit under fixed continuation policies

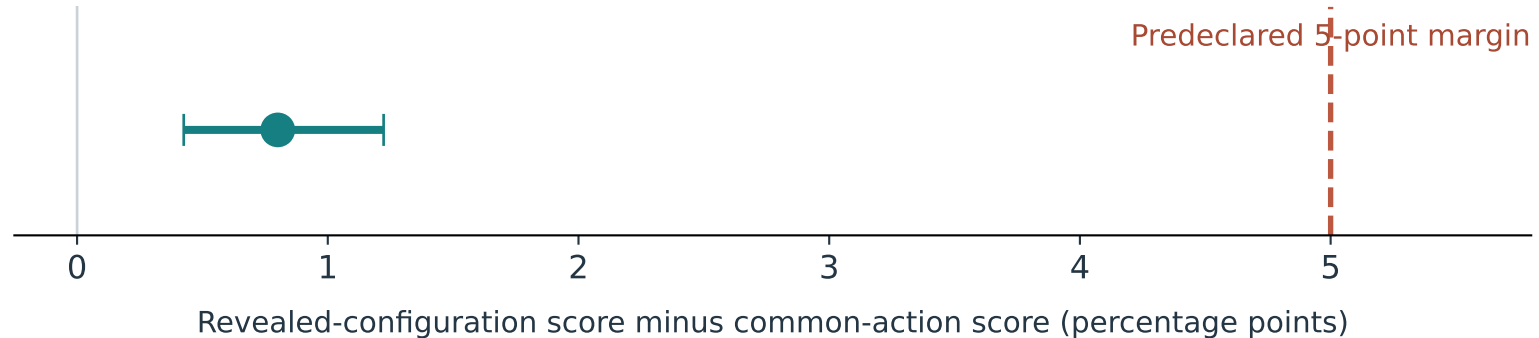
Fixed primary cohort: 300 positions from 300 games with $H < 15$

+0.80

percentage points

Two-sided 95% game-bootstrap CI: +0.43 to +1.22

One-sided 95% upper: 1.16 < 5 percentage points



What happened across all tested actions and continuations?



All won: 183/300 | All lost: 22/300 | Results varied: 95/300

Sampled outcomes, not proved forced wins or losses.

Tie-aware reference stability: 98.3%. All planned primary positions passed data checks.

Finite sampled configurations and fixed continuation policies; both selectors know the decklists.

H counts opponent deck + hand + Prizes. These results do not bound an optimal perfect-information solver.